

Image Classification With EMNIST - An Extended Version of MNIST

Tamir Bennatan

tamir.bennatan@mail.mcgill.ca
260614526

Kyle Levy

kyle.levy@mail.mcgill.ca
260612028

Phillip Ryjanovsky

Philip.ryjanovsky@mail.mcgill.ca
260604024

I. INTRODUCTION

The MNIST dataset - comprising of images of handwritten digits of 10 classes - is a popular dataset, widely used as a benchmark for learning and classification tasks amongst the Neural Network and Computer Vision communities. MNIST enjoys several favorable characteristics which have contributed to its widespread adoption (Cohen, et. al; 2017). First, the dataset is relatively small (each of the 60,000 examples are (28x28) pixel grayscale images), which makes it feasible to work with the dataset on most processors [Figure 1]. Second, the digits have been centered within each image, have consistent heights and similar widths. The digits are all distinguishable to the human eye, and the labels associated with each image are accurate. Because of these properties, many researchers have achieved near perfect accuracy in the image classification task on MNIST (for example: Lecun et. al; 2001).

In this project, we address the digit classification task on a more difficult dataset called Extended-MNIST (EMNIST). Each example in the EMNIST dataset was generated with the following procedure:

- 1) Randomly select 2 or 3 images from MNIST.
- 2) Rescale each image by a factor ranging from 40% to 120% the original size (the label of the generated image is that of the MNIST image which was scaled up by the largest factor).
- 3) Apply random rotations to each image.
- 4) Superimpose these rotated images onto a random background.

The inherent randomness in the EMNIST examples makes image classification on EMNIST more difficult than on MNIST [Figure 2].

In this report, we discuss the pre-processing measures we took to address the noisy features of the EMNIST examples. We then summarize our process in fitting and tuning Logistic Regression, Feedforward Neural Network, and Convolutional Neural Network (CNN) classifiers. CNNs proved most effective; we achieved 92.6% validation accuracy using the aggregated predictions of four CNN models. We then summarize our results and our shortcomings, and ways in which we can improve our accuracy in the future.

II. FEATURE DESIGN

The EMNIST dataset introduces several difficulties in building an effective classifier. The random backgrounds of each example in the EMNIST dataset were generated independently of the examples' labels. Thus, they only add

noise to each image, which can impede effective learning by a machine learning algorithm.

It is also unclear which digit in each image was scaled up the most during the construction of each example in EMNIST, as we do not have prior knowledge of the dimension of the MNIST digits used to create each example in EMNIST. This makes it difficult to isolate the digit within each EMNIST example that corresponds with the target label.

Finally, an effective classifier for the EMNIST dataset must be robust to rotations in images, as the EMNIST dataset was built by rotating MNIST images randomly. As most machine learning algorithms do not encode images in a way that is robust to rotations, this adds difficulty to the classification task.

A. Removing Noisy Backgrounds

The first preprocessing step we took was to remove the backgrounds of each example in EMNIST. Luckily, the pixels in each example corresponding to a digit extracted from MNIST have a fixed integer value of 255 (corresponding to the color white). This made it possible to replace the random backgrounds with black backgrounds, by converting all non-white pixels to black pixel values [Figure 3].




Fig. 3. Three sample images from EMNIST, before filling the backgrounds (top) and after filling the backgrounds (bottom).

B. Extracting the 'Largest' Digit

Each example in EMNIST contains several digits, but only one label. Intuitively, this makes the learning task difficult, as a machine learning algorithm will try to associate the properties of several digits with an example's label, when in fact the task is to only correctly identify the *largest* digit.

Thus, we set out to extract the "*largest*" digit from each example in EMNIST. Since we do not know the original dimensions of each of the digits in each image, we employed a series of heuristics to define the size of a digit. These heuristics are:



Fig. 1. Three sample images from the MNIST dataset.



Fig. 2. Three sample images from the EMNIST dataset.

- 1) **Bounding Rectangle Area:** The area of the smallest upright rectangle that circumscribes the digit.
- 2) **Bounding Circle Area:** The area of the smallest circle which escribes the digit.
- 3) **Minimum Rotated Rectangle Area:** The area of the smallest (possibly rotated) rectangle that escribes the digit.
- 4) **Minimum Rotated Rectangle Maximum Dimension**
The largest dimension (height,width) of the smallest (possibly rotated) rectangle that circumscribes the digit.

We refer to these huristics as *heuristics 1-4*, respectively.




Fig. 4. A sample image from EMNIST, superimposed by the geometric objects used to measure the size of each image using heuristics 1-4. Heuristic 1 is the top left, Hueristic 2 is the top right, Heuristic 3 is the bottom left, and Heuristic 4 is the bottom right.

We treated the choice of which heuristic to use to measure the size of a digit as a hyperparameter, selected from them using cross-validation. Once we selected a heuristic, cropped each example to the a square which circumscribes the largest digit (according to this heuristic), and resized this square to 28 by 28 pixels.

C. Data Augmentation

To help make our model more robust to rotated digits, we augmented our training data with randomly rotated copies

of each training example. Our rational is that if we train on randomly rotated images, our classifier will scale unseen randomly rotated images more effeciely. We tuned the number of times we rotated each image (called the "*batch size*"), and the range of degrees with which we rotated each image as hyperparameters.

III. ALGORITHMS

We trained classifiers of three types: Regularized Logistic Regression, Feedforward Neural Network (FFNN), and Convolutional Neural Networks (CNN). We trained each of these classifiers on the datasets processed using each of size heuristics described above, though we only experimented with data augmentation when training the CNN.

Regularized Logisitic regression served as our benchmark, as it is a simple and easy to implement algorithm.

For all experiments using FFNNs, we learned the parameters using stochastic gradient descent. We used the sigmoid activation function for all hidden units and the output units, and took a model's predictions for example X_i to be the class which has the highest output activation:

$$\hat{y}_i = \underset{\text{Output Units; } a_i}{\arg \max} \sigma(a_i)$$

Where σ is the logistic sigmoid function. Thus, the loss function we used was the multi-class log-loss, which is defined for a k class classification problem as:

$$Loss(\hat{y}_i, y_i) = -1 \sum_{k=1}^K * (y_{i,k} \log(\hat{y}_{i,k}) + (1 - y_{i,k}) \log(1 - \hat{y}_{i,k}))$$

Where $y_i \in \mathbb{R}^k$ is a one-hot representation of the i^{th} output, and $\hat{y}_{i,k}$ is the predicted probability that example i belongs in class k .

Finally, we trained CNNs of four different archetecutres - all inspired by the LeNet architecture by Lecun et. al (Legun et. al; 1995). These architectures alternate between convolutional and max-pooling layers, and output predictions through two fully connected layers. One archetecutre had three convolutional layers, two had four convolutional layers, and one had six convolutional layers.




Fig. 5. One of the four CNN architectures we experimented with. The other three differ slightly in the kernel size, pooling size, and number of convolutional layers. Image modified from [4].

A summary of the four architectures we experimented with is presented in the appendix. We refer to them as *Architecture 1-4*.

IV. METHODOLOGY

We started by splitting the labeled data into training and validation splits of sizes 40,000 and 10,000 examples, respectively. We kept these splits consistent in all experiments, so that we could gauge the relative effectiveness of the models. For each of the three algorithms we used, we chose which size heuristic(s) to use based on which heuristic(s) yielded the highest validation accuracy when used to select the largest digit in each image during preprocessing. We also used this validation set to tune model-specific parameters.

A. Regularized Logistic Regression: Model Hyperparameters

The two hyperparameters we tuned for logistic regression were the inverse-regularization coefficient, C^1 , and the regularization loss function (either $l1$ or $l2$). Using the processed datasets generated from using each of the four size heuristics, we ran a grid-search over 24 hyperparameter combinations, and selected the combination that yielded the best validation accuracy.

B. Feedforward Neural Network: Model Hyperparameters

The two hyperparameters we tuned when using FFNNs were the learning rate, α , and the model architecture. We did not implement regularization.

We first trained with a learning rate of $\eta = .01$, but observed that the loss was not decreasing monotonically. We decreased η to $.001$, and observed smoother convergence. Thus, we fixed the learning rate to $\eta = .001$ for all further experiments.

C. Feedforward Neural Network: Learning Optimizations

We began by trying to fit our models with simple unregularized stochastic gradient descent (SGD), but quickly found that our models were not converging.

To address this, we implemented *momentum* to help calculate the amount each weight should be updated on iteration t , denoted ν_t , where ν_t is defined:

$$\nu_t = \gamma \nu_{t-1} + \eta \nabla_{\theta} \text{Loss}(\theta)$$

Where γ is an additional hyperparameter introduced.

¹In the notation seen in class, $C \propto 1/\lambda$.

TABLE I
HYPERPARAMETER RANGES EXPLORED FOR CNN MODELS

Hyperparameter	Range Explored
Model Architecture	{Architecture 1, Architecture 2, Architecture 3, Architecture 4}
Data Augmentation - (Batch size, Rotation Range)	{(1, [0, 0]), (8, [-10, 10]), (16, [-20, 20]), (16, [-30, 30])}
Learning Rate	{.000001, .00005, .0001}
Digit Size Heuristic	{Heuristic 1, Heuristic 2, Heuristic 3, Heuristic 4}

We also implemented *power scaling* - a method of slowly reducing the learning rate η with each epoch. At the end of each epoch, η is updated as follows:

$$\eta_{\text{new}} = \frac{\eta_{\text{old}}}{(\text{Epoch number})^p}$$

For a scaling parameter p . This allows the weights to change drastically in the beginning of training, and slowly dampens the change in weights as training progresses, so that local maxima are not overlooked.

We found that introducing these two learning optimizations helped the SGD algorithm minimize loss much faster. Thus, for all experiments with FFNN, we used both momentum and power scaling, with the hyperparameters fixed at $\gamma = .9, p = .5$.

D. Convolutional Neural Network: Hyperparameter Tuning

We considered many hyperparameters when tuning our CNN models. The most important hyperparameters are summarized in Table 1:

We started by training all four architectures on the data preprocessed using Heuristic 1 with the high learning rate of $.0001$. After observing that architectures 2 and 4 yielded much higher validation accuracy than architectures 1 and 3, we decided to omit the latter architectures from all further experiments. We also decided to limit the range of the learning rate η to $\eta \in \{.000001, .00005\}$, as we saw that convergence was occurring with high variance.

We then took a random sample of 64 images from EMNIST, and extracted the "largest" digit from each image using each of the four size heuristics defined. We observed by visual inspection that when using Heuristics 3 and 4, we isolated the digits corresponding to the labels more frequently than when we used Heuristics 1 and 2. We decided to use data preprocessed with Heuristics 3 and 4 in all further experiments.

We then performed a gridsearch on the remaining hyperparameter combinations permitted by the constrained hyperparameter ranges [Table 2], totalling parameter combinations.

E. CNN: Learning Optimization

We learned the weights of our networks using the *RM-SProp* algorithm, with the hyperparameter ρ fixed at 0.9. While learning, we used an early stopping scheme, which caused training to stop if the validation accuracy has not improved for 8 consecutive epochs.

TABLE II
CONSTRAINED HYPERPARAMETER RANGES

Hyperparameter	Range Explored
Model Architecture	{Architecture 2, Architecture 4}
Data Augmentation - (Batch size, Rotation Range)	{(1, [0, 0]), (8, [-10, 10]), (16, [-20, 20]), (16, [-30, 30])}
Learning Rate	{.000001, .00005}
Digit Size Heuristic	{Heuristic 3, Heuristic 4}




Fig. 6. Learning curves for training periods with $\eta = .0001$ and $\eta = .00005$, with early stopping. Note: training accuracy is lower than validation accuracy. This is because we used data augmentation, and the new randomly rotated images made the training data “harder” to learn than the validation data.

At first, we tried learning using learning rate of $\eta = .0001$, but we observed through the use of learning curves that the convergence had very high variance. After reducing the learning rate to $.00005$, we achieved smoother learning [Figure 6].

F. CNN: Aggregated Predictions

We initialized the weights of our CNN’s randomly. As such, every time we fit a model with a certain configuration, the final weights may be different, leading to different predictions.

To account for this randomness, we fit two models of each hyperparameter configuration, each with different random initialization. This allowed us to keep the better of the two trained models (based on validation-set accuracy) for each hyperparameter configuration.

Thus, for every combination of the hyperparameters (Data Augmentation Strategy, Size Heuristic), we had two models fit using Architecture 2, and two models fit using Architec-

ture 4. This gave us the opportunity to aggregate predictions² across several trained models.

Two of the aggregations we experimented with increased validation accuracy by over 1%, which we considered substantial. We refer to these aggregations as *Aggregation 1* and *Aggregation 2*.

Aggregation 1:

- Predictions of four models aggregated.
- Two configurations - each repeated once with a different random initialization.
- Configuration 1: Architecture 2, Batch Size = 16, Rotation Range [-.2, .2], $\eta = .00005$, Size Heuristic 3
- Configuration 2: Architecture 4, Batch Size = 16, Rotation Range [-.2, .2], $\eta = .00005$, Size Heuristic 4
- Configuration 3: Architecture 2, Batch Size = 16, Rotation Range [-.2, .2], $\eta = .00005$, Size Heuristic 3
- Configuration 4: Architecture 4, Batch Size = 16, Rotation Range [-.2, .2], $\eta = .00005$, Size Heuristic 4

Aggregation 2:

- The same four configurations as in Aggregation 1, but each model configuration was trained and included twice.

V. RESULTS

In this section we discuss the hyperparameters which had the largest impact on the performance of each model, and then we compare the performance of each of our tuned models on the validation set.

A. Regularized Logistic Regression

After doing a thorough search through many (C, Regularization Loss Function) combination, we found that the hyperparameter which had the largest impact on validation accuracy was the regularization loss function. The best model trained with $l1$ -loss achieved a validation accuracy over 1% higher than the best model trained with $l2$ -loss [Table 3].

TABLE III
BEST LOGISTIC REGRESSION MODELS, TRAINED WITH L1/L2 LOSS

Loss Function	Regularization Parameter	Validation Accuracy
l1-loss	0.09	0.731
l1-loss	.27	0.735
l2-loss	.01	0.723
l2-loss	.03	0.721

Models trained with $l1$ -loss regularization find sparse solutions - meaning that many input features do not contribute to the decision boundary. This has an interesting interpretation for the EMNIST classification problem; it means that the Logistic Regression model learned that some pixel values (features) do not help to predict the largest digit, and should be ignored.

²To aggregate predictions of several models, we used the class which was predicted the most frequently amongst each individual model to be the aggregated class prediction.

For example, since we padded all images with a black border during preprocessing, these pixel values do not help to distinguish between the digits. Its reasonable to assume that the Logistic Regression model, trained with $l1$ -loss learned to ignore the values of those pixels.

B. Feedforward Neural Network

We experimented with three model architectures when fitting our FFNN:

TABLE IV
FFNN ARCHITECTURES

Architecture Number	Number of Hidden Layers	Hidden Layer Dimensions
Architecture 1	1	(128)
Architecture 2	2	(64, 20)
Architecture 3	2	(128, 64)

We trained all models for 20 epochs with SGD, using both momentum and power scaling.

After some experimentations, we found that the configuration $\eta = .01, \gamma = .9, p = .9$ led to a reasonable learning rate. Using these hyperparameter configurations, we fit each of the architectures described in table 3, and got the following results:

TABLE V
VALIDATION ACCURACIES FOR FFNN EXPERIMENTS

Architecture Number	Validation Accuracy
Architecture 1	.55
Architecture 2	.64
Architecture 3	.64

VI. DISCUSSION AND CONCLUSIONS

Using the three sets of features described in this paper, we were able to improve the accuracy of our baseline model from 74.56% to 79.25%. Of the three feature sets, those computed using Tf-Idf scores were the most impactful. This confirms our hypothesis that measures of word importance and specificity are useful in comparing the semantic similarity between two questions.

We also applied and extended a framework for modelling semantic similarities between short texts using knowledge-based word similarity scores, proposed by Mihalcea et al (2006). We found that these scores were useful in our classification task. These features were highly correlated with those calculated using Tf-Idf scores, so it may be that they are merely capturing the specificity of the words in each question, rather than the semantic similarity of these words.

Finally, we found that our answer type predictions were not useful for detecting duplicate questions. We suspect that this is because we used overly complex models to predict the answer type of a question. The TREC dataset has only 6,000 observations, but the number of parameters in our LSTM classifiers ranged from 70,506 to 28,141,268. It's probable that these classifiers were overfit to the TREC dataset, and did not generalize to the Quora dataset.

In future work, we would like to try using simpler models for predicting a question's answer type - such as Naïve Bayes or a Support Vector Machine. We could also try to incorporate predictions of the fine answer categories in the annotated TREC dataset, as apposed to the coarse categories we used.

We also think that an interesting extension to this work would be to extract features that concern the grammatical functions of the words in each question. For example, one could construct a dependency parse tree for each question, and compare the subjects and direct objects in each question, as well as the head words of their repective parses.

VII. STATEMENT OF CONTRIBUTION

Tamir engineered the features and classifiers described, worked on the paper, and conducted model experiments. Jack contributed to the classifiers and paper, and experimented with advanced techniques for feature extraction, such as named entity recognition and headword extraction.

VIII. APPENDIX

REFERENCES

- [1] Cohen, G., Afshar, S., Tapson, J., & van Schaik, A. (2017). EMNIST: an extension of MNIST to handwritten letters. Retrieved from <http://arxiv.org/abs/1702.05373> (link is external)
- [2] Y. LeCun, L. Bottou, Y. Bengio and P. Haffner: Gradient-Based Learning Applied to Document Recognition, Intelligent Signal Processing, 306-351, IEEE Press, 2001,
- [3] Y. LeCun, L. D. Jackel, L. Bottou, A. Brunot, C. Cortes, J. S. Denker, H. Drucker, I. Guyon, U. A. Muller, E. Sackinger, P. Simard and V. Vapnik: Comparison of learning algorithms for handwritten digit recognition, in Fogelman, F. and Gallinari, P. (Eds), International Conference on Artificial Neural Networks, 53-60, EC2 & Cie, Paris, 1995.

BLOG POSTS

- [4] Theart, R. (2017, November 29). Getting started with PyTorch for Deep Learning (Part 3: Neural Network basics). Retrieved from <https://codetolight.wordpress.com/2017/11/29/getting-started-with-pytorch-for-deep-learning-part-3-neural-network-basics/>

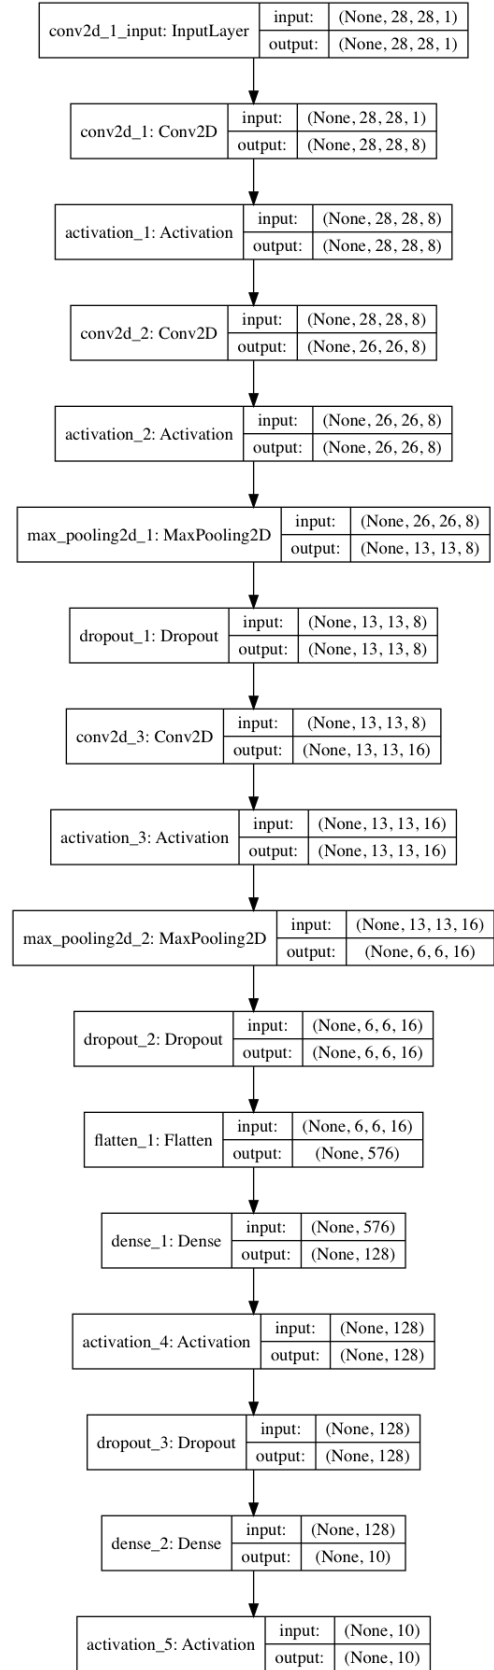


Fig. 7. Architecture 1: Three convolutional layers, 76,978 trainable parameters.

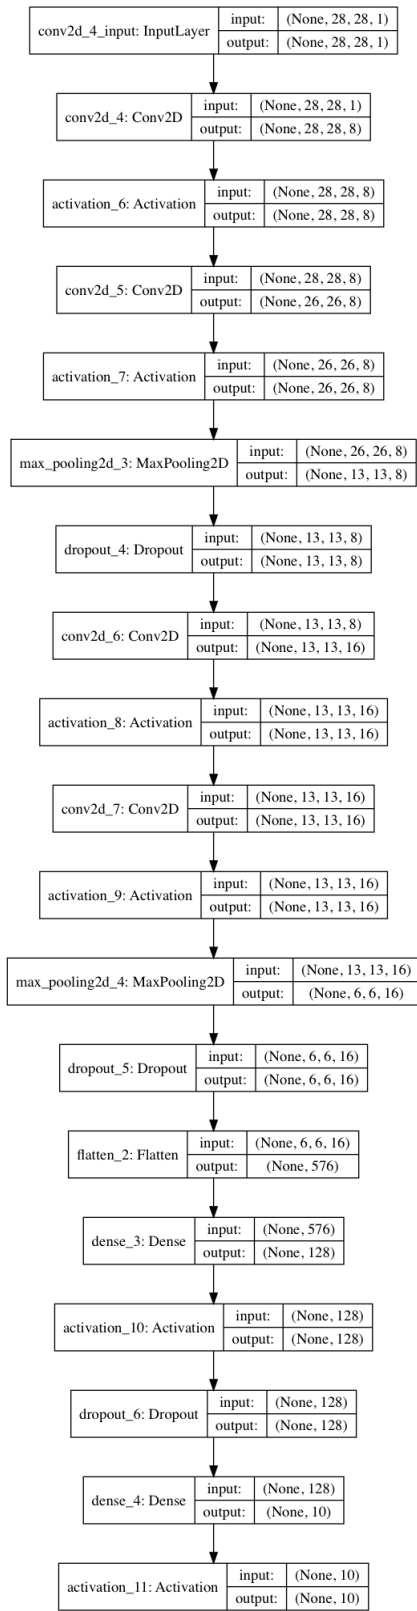


Fig. 8. Architecture 2: Four convolutional layers, 79,298 trainable parameters.

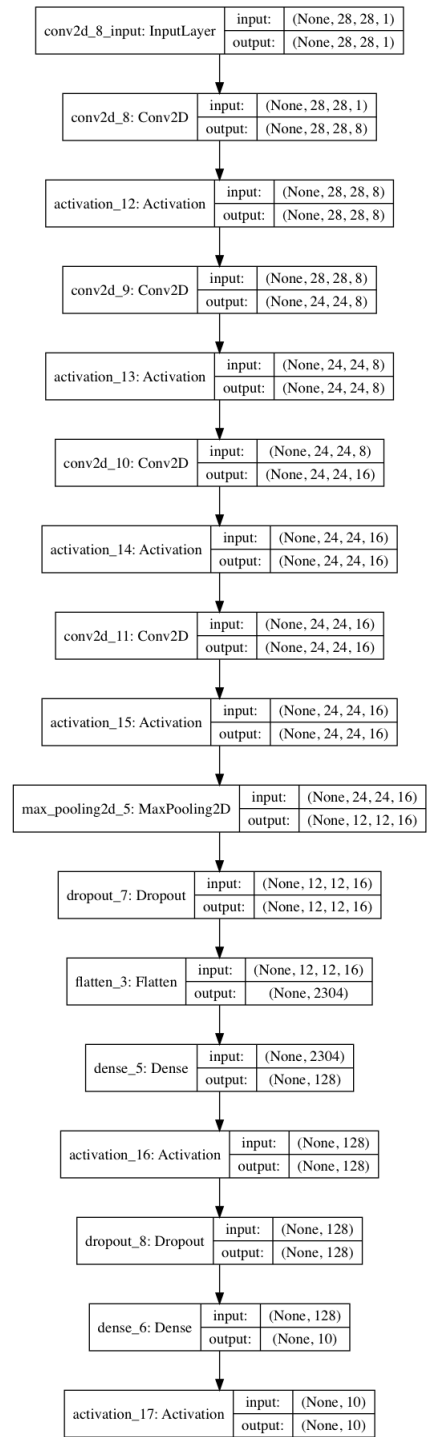


Fig. 9. Architecture 3: Four convolutional layers, 307,778 trainable parameters.

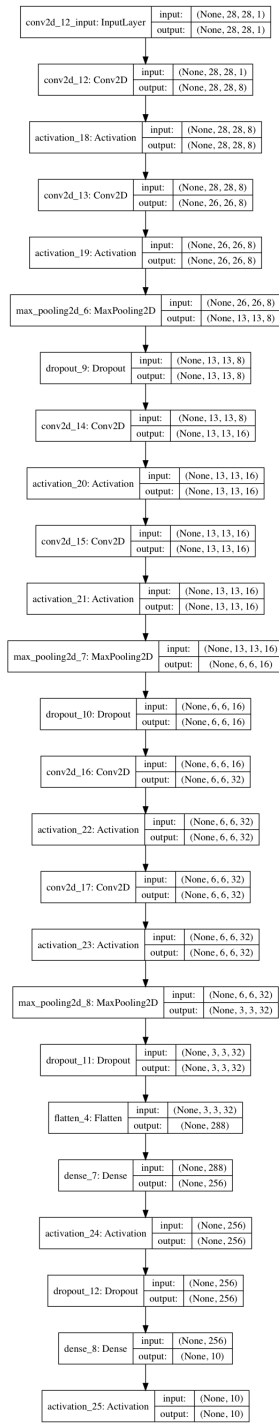


Fig. 10. Architecture 4: Six convolutional layers, 94,594 trainable parameters.